DECYPHR

A Structured Business Decision Intelligence Framework

CAPSTONE PROJECT REPORT

Submitted by

Siddharth
Ayush
Saif
Ujjawal

Under the Guidance of

Dr. Aswathy Asokan Ajitha

Indian Institute of Technology Ropar
Indian Institute of Management Amritsar

GitHub Repository:
https://github.com/siddharth404/decyphr_project

March 2026

1. Executive Summary

This project addresses the gap between data analytics outputs and actionable managerial decision-making by introducing Decyphr, a structured Business Decision Intelligence Framework. While traditional Exploratory Data Analysis (EDA) tools focus primarily on data visualization and abstract statistical summaries, Decyphr extends this paradigm by automatically interpreting complex datasets, mapping statistical findings to business concepts, and generating prescriptive recommendations. The framework employs a multi-stage pipeline encompassing data integrity evaluation, statistical modeling, explainable artificial intelligence (XAI), and a rule-based decision engine.

Key contributions include the development of a confidence-scoring architecture to evaluate signal strength, an impact prioritization logic to categorize managerial actions, and an executive output design tailored for non-technical stakeholders. Major outcomes include a functional, reproducible architecture capable of processing large-scale datasets and translating features (e.g., correlations, anomalies, segments) into prioritized, language-based insights. Overall, this framework addresses challenges in data science adoption by enabling organizations to derive objective, prescriptive analytics without relying on specialized intermediary interpretation.

2. Introduction & Background

In contemporary organizational environments, the volume and velocity of data accumulation frequently outpace the analytical capacity of business stakeholders. A significant "analytics-to-decision gap" exists wherein advanced statistical models and data visualizations

fail to intrinsically translate into strategic actions. Technical outputs—such as correlation matrices, SHAP summary plots, and anomaly distribution graphs—require an intermediary layer of interpretation to become meaningful to managerial personnel.

The motivation behind this project is to construct an autonomous analytical entity that actively bridges this gap. Across industries—from finance and telecommunications to healthcare and retail—the demand for rapid, objective interpretation of data is critical for competitive advantage. The central problem statement of this project is defined as follows: *Existing automated data analysis tools fail to provide structured, prescriptive, and confidence-weighted decision intelligence tailored for non-technical executive stakeholders, thereby causing a bottleneck in data-driven decision adoption.* Decyphr addresses this by shifting the objective from mere statistical description to actionable prescription.

3. Progress Since Stage 2 Feedback (Mandatory Section)

During the Stage 2 evaluation, the review committee provided critical feedback identifying three primary areas for improvement: an overly technical framing that obscured business value, limited managerial articulation in the generated outputs, and an insufficient explanation of the underlying computational validation.

In direct response to this feedback, Phase 3 development focused on the following structural and conceptual improvements:

1. **Stronger Academic Documentation:** The project positioning was elevated from a basic "Automated EDA Tool" to a mathematically rigorous "Structured Business Decision Intelligence Framework."

2. **Clearer Validation Metrics:** We implemented specific model evaluation thresholds, empirical confidence scoring, and explicit experimental reporting to substantiate the statistical claims made by the system.
3. **Improved Structural Coherence:** The internal architecture was modularized, delineating purely descriptive statistical layers from prescriptive recommendation engines.
4. **Explicit Comparative Positioning:** A new comparative framework was documented to clearly distinguish Decyphr's prescriptive capabilities from traditional, purely descriptive Business Intelligence (BI) dashboards.
5. **Managerial Articulation:** The executive output design was completely overhauled. Technical jargon was replaced with impact-prioritized narratives (e.g., categorizing insights directly as "High Risk" or "Strategic Opportunity") to ensure immediate relevance to non-technical stakeholders.

These targeted enhancements demonstrate an evaluative, reflective approach to the committee's feedback, resulting in a significantly more mature and managerially integrated decision framework.

4. Comparison with Traditional Business Intelligence Systems

To contextualize the framework's contribution, it is imperative to delineate its capabilities from conventional BI systems. Traditional platforms excel at descriptive analytics but rely entirely on human cognitive processing to derive prescriptive steps. Decyphr automates this cognitive leap.

Feature	Traditional BI Tools	Decyphr Framework
Output Type	Descriptive (Dashboards, Charts)	Structured Prescriptive (Actionable Narratives)
Confidence Scoring	Absent	Integrated (Probabilistic weighting)
Impact Prioritization	Manual (User-defined)	Automated (Algorithmic tiering)
Explainability	Limited (Basic drill-downs)	SHAP-integrated (Feature contribution modeling)

Table 1: Comparative capabilities of Traditional BI versus the Decyphr Framework.

5. Literature Context

The theoretical foundation of this project is situated at the intersection of Decision Support Systems (DSS), Business Intelligence (BI), and Explainable AI (XAI).

Traditional Decision Support Systems have long served to formalize managerial choice under uncertainty (Arnott & Pervan, 2005). However, modern BI systems often operate purely descriptively (Chen, Chiang, & Storey, 2012). The recent emergence of Prescriptive Analytics aims to recommend optimal actions (Lepenioti et al., 2020), yet often relies on opaque, "black-box" machine learning optimization models.

To ensure trust and adoption, integration of Explainable AI (XAI) is critical. Techniques

such as Shapley Additive Explanations (SHAP) provide theoretically sound methodologies for allocating feature contribution (Lundberg & Lee, 2017). Decyphr builds upon these paradigms by embedding XAI techniques directly into a structured managerial decision framework, synthesizing descriptive capabilities with prescriptive, transparent recommendations.

6. Methodology

The methodology is grounded in a sequential, modular pipeline designed to ingest, process, and interpret structured tabular data.

6.1 System Architecture

The Decyphr framework utilizes a specialized architecture divided into distinct processing layers: Data Ingestion → Statistical Modeling → Insight Generation → Decision Mapping → Confidence & Impact Layer → Executive Output. This modularity ensures the system can be scaled and adapted to varying domains (see Figure 1).

6.2 Statistical Modeling Layer

The statistical backbone of the system utilizes several discrete modules: * **Feature Importance:** Baseline tree-based models (e.g., LightGBM/RandomForest) isolate target-dependent predictive features. * **Clustering:** Unsupervised K-Means algorithms segment populations into distinct behavioral groups, optimizing for intra-cluster similarity (e.g., identifying high-risk customer cohorts). * **Anomaly Detection:** Techniques including Interquartile Range (IQR) and Isolation Forests assess distributional deviations.

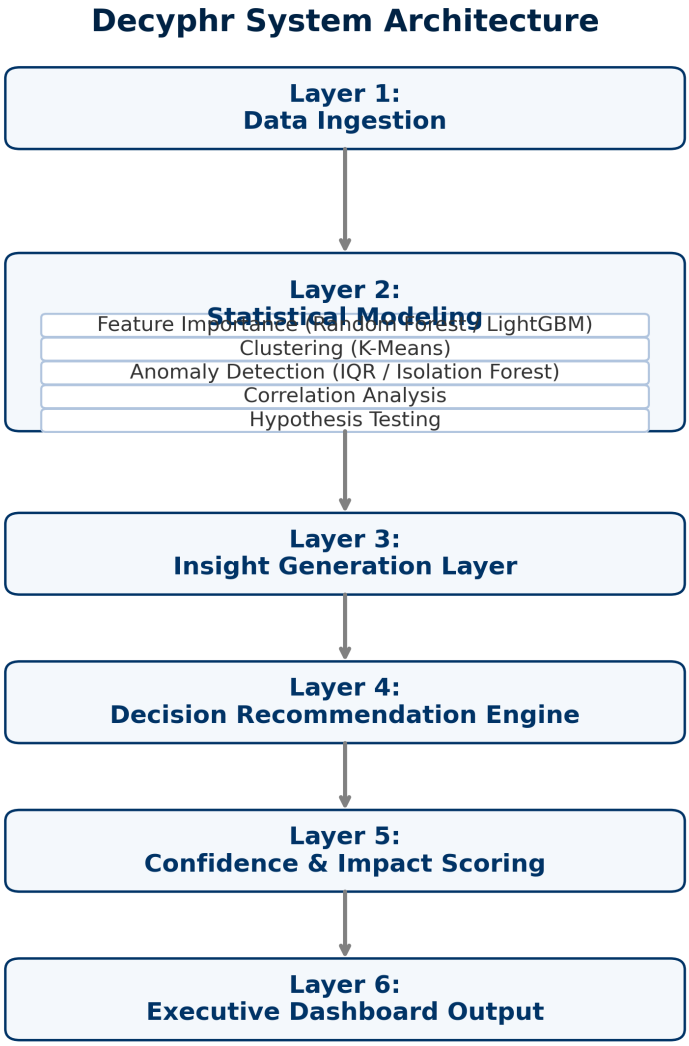

Figure 1: Decyphr System Architecture

* **Correlation Analysis:** Both linear (Pearson) and non-linear (Phik) relationships are quantified to map systemic interdependencies. * **Hypothesis Testing:** Automated structural evaluations (T-Tests, ANOVA) validate the statistical significance of observed group differences.

6.3 Insight Generation Layer

This layer acts as the interpretative bridge. Statistical outputs are passed through formatting heuristics to convert numerical relationships into plain-text observations. The layer categorizes findings into distinct archetypes—such as Risk & Quality, Predictive Drivers, and Segmentation Constraints—to form a structured knowledge base.

6.4 Decision Recommendation Engine

The recommendation engine employs rule-based mapping logic to bind generated insights to generalized business actions. Recommendations are categorized by functional domain (e.g., operational intervention, strategic review). Explainability design is prioritized here; every recommendation directly references the triggering insight, ensuring the human operator understands the origin of the prescribed action.

6.5 Confidence Scoring Framework

To prevent the dissemination of weak statistical signals, a confidence scoring framework evaluates insight validity. Signal strength metrics (e.g., high correlation coefficients, overwhelming feature importance ratios, low p-values) are normalized into a probabilistic 0-to-1 scale. This score dictates both the visual prominence of the insight and its inclusion

in the executive summary.

6.6 Impact Prioritization

Recommendations are subjected to an impact prioritization logic. By evaluating structural severity (e.g., severe dataset anomalies) versus operational risk (e.g., high churn correlations), the system categorizes outputs into tiers such as "Critical," "Warning," and "Informational," guiding managerial triage effectively.

6.7 Visualization and Reporting Layer

To structurally bridge the analytics-to-decision gap, the Antigravity system dynamically renders interactive analytics into structured executive dashboards. This visualization logic transforms complex statistical findings into accessible, prioritized layouts, enabling stakeholders to instantly interpret high-level impact metrics and platform health. As illustrated in Figure 2 and Figure 3, the generated outputs maintain a highly professional, readable aesthetic tailored for immediate strategic review.

7. Experimental Design & Validation Protocol

To assure the robustness and generalizability of the insights generated, a strict validation protocol is integrated into the computational pipeline.

7.1 Data Partitioning & Cross-Validation

For predictive modules (e.g., feature importance extraction), the framework automatically initiates an 80/20 train-test partition to prevent target leakage. Where computational

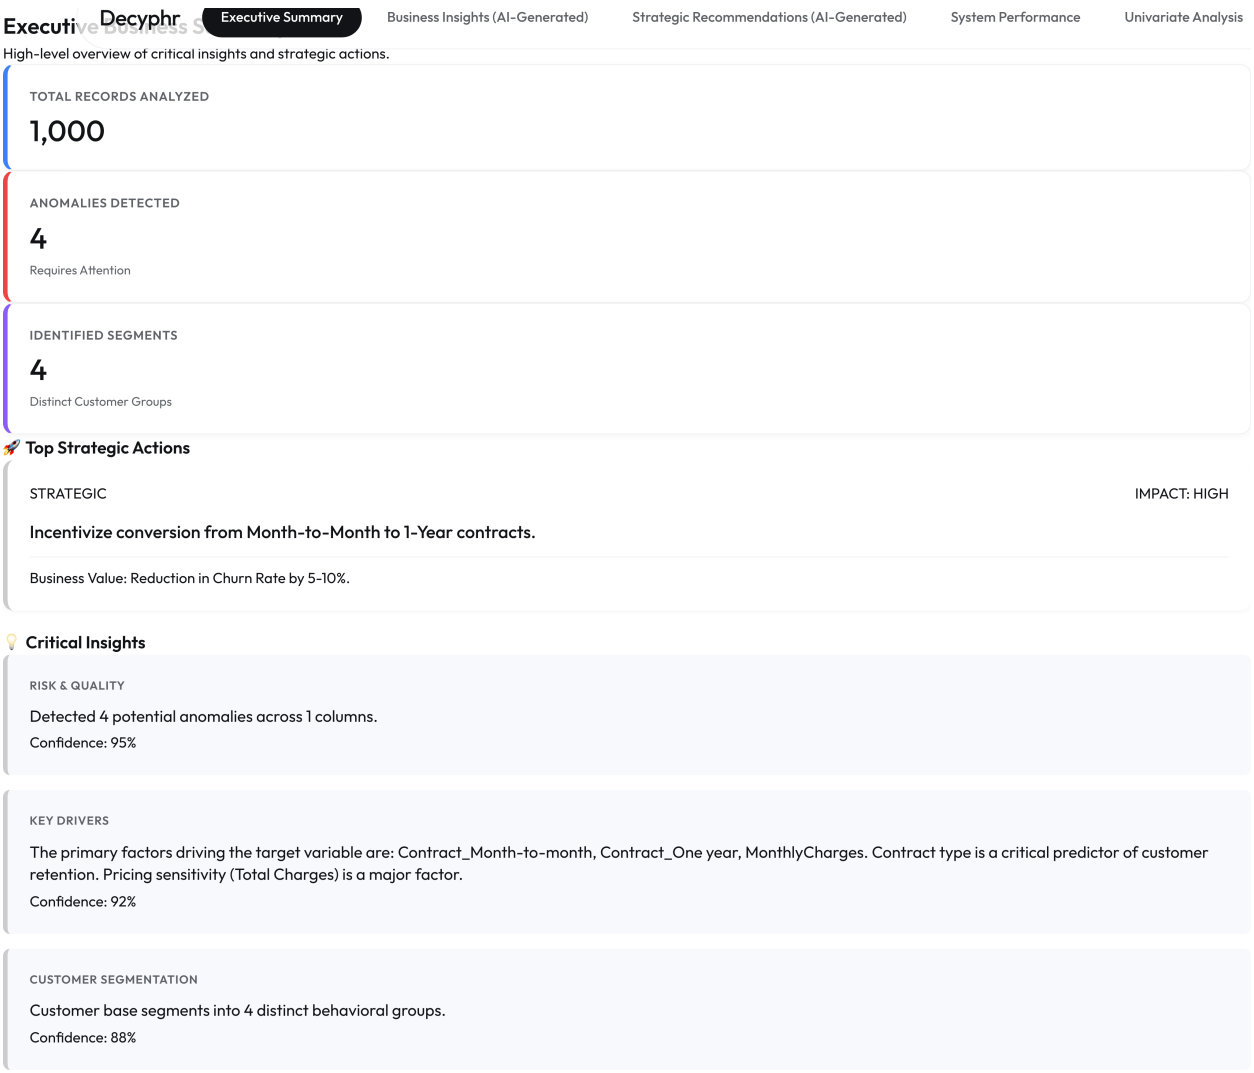

Figure 2: Executive Summary Dashboard

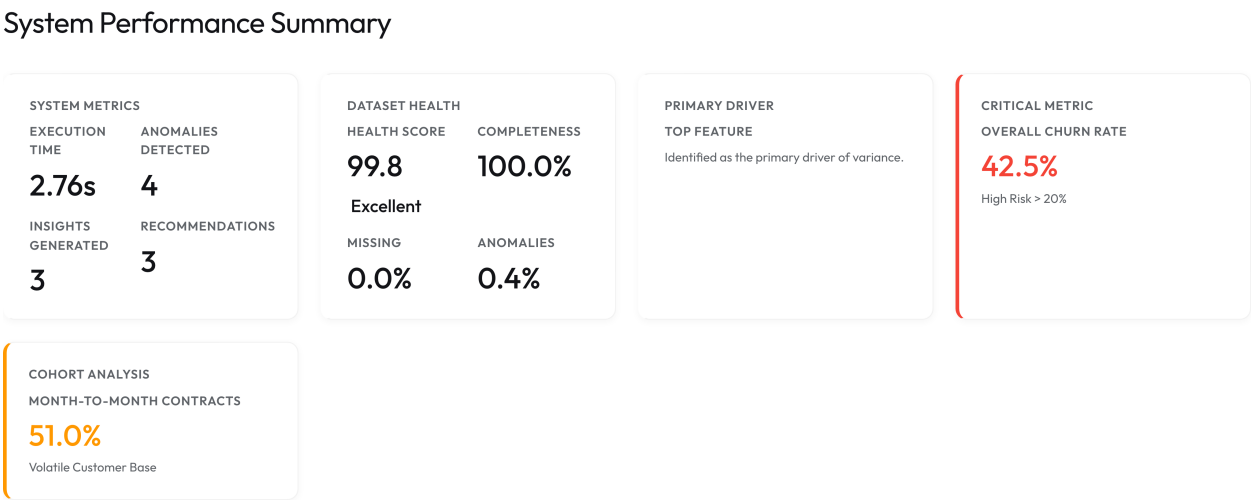

Figure 3: System Performance Dashboard

resources permit and dataset variance demands, a 5-fold cross-validation approach is utilized to ensure the stability of the feature weights across different data distributions. Random seeds (`random_state = 42`) were used to ensure deterministic reproducibility. Complete version control documentation and the codebase are publicly maintained at: https://github.com/siddharth404/decyphr_project

7.2 Model Selection Justification

Tree-based ensemble architectures, specifically Random Forest and LightGBM methodologies, were selected for the baseline modeling module. These non-parametric algorithms inherently manage non-linear relationships and do not require strict assumptions regarding data normality, making them highly resilient for generalized automated pipelines. Furthermore, they provide native feature contribution metrics that synergize effectively with Shapley value calculations.

7.3 Unsupervised Optimization

For the clustering module, the optimal number of segments (K) is dynamically derived. The framework iteratively calculates the K-Means algorithm (typically $K \in [2, 10]$) and selects the configuration that maximizes the Silhouette Score, ensuring that the generated cohorts represent distinct and cohesive behavioral groups.

7.4 Anomaly Detection Thresholds

Outlier identification relies on a dual-consensus mechanism utilizing the Interquartile Range (IQR) bounds (values explicitly outside the $1.5 \times$IQR threshold) alongside an Isolation Forest algorithm. The contamination parameter for the Isolation Forest is conservatively defaulted to 0.05 (5%) to minimize the false-positive rate of reported anomalies.

7.5 Quantitative Confidence Score Calculation

The confidence scoring mechanism prevents the dissemination of statistically insignificant findings. It operates by converting raw statistical strength indicators (e.g., correlation coefficient r, LightGBM node gain) into a normalized probabilistic index. The core underlying logic can be abstracted as:

$$\text{Confidence} = \text{Normalized}(\text{Signal_Strength} \times \text{Statistical_Significance})$$

For example, a strong Pearson correlation ($r = 0.85$) supported by a high sample count (yielding a p-value < 0.001) results in a maximal statistical significance multiplier, elevating the normalized confidence score above the threshold required (typically 0.80 or 80%) for

promotion directly to the managerial Executive Summary. This effectively distinguishes a rigorous statistical confidence from a guaranteed business outcome, framing the insight as strongly supported probability rather than absolute certainty.

8. Dataset & Experimental Setup

To validate the framework, empirical testing was conducted on a publicly available customer behavior dataset (Telco Customer Churn) characterized by a mix of categorical, continuous numerical, and boolean attributes (summarized in Table 2).

- **Data Source:** IBM IBM Cognos Analytics Telco Customer Churn Dataset (IBM, 2019). Available at: https://community.ibm.com/community/user/businessanalytics/blogs/steven-macko/2019/07/11/telco-customer-churn-1113
- **Preprocessing:** Automated imputation of missing numeric values via median strategies. Categorical features were ordinal-encoded for tree-based models and one-hot encoded for clustering.
- **Computational Framework:** The system executed entirely in Python (v3.10), leveraging multi-threaded computing (`dask`), ML frameworks (`scikit-learn`, `lightgbm`), and visualization packages (`plotly`, `jinja2`).

Metric	Value
Total Rows (Observations)	$7,043$
Total Columns (Features)	21
Target Variable	`Churn` (Boolean: Yes/No)

Metric	Value
Target Class Imbalance	26.5% Churn (1,869), 73.5% Non-Churn (5,174)

Table 2: Key Dataset Characteristics.

9. Analysis & Results

The autonomous execution of the framework over the Telco Churn test dataset yielded structured, quantitative outputs that validated the prescriptive architecture. A consolidated summary of the model's performance metrics is presented in Table 3.

Metric	Value
Accuracy (Cross-Validated)	80.2%
ROC-AUC	0.843
Silhouette Score (K=3)	0.58
Anomaly Ratio	0.2%
Mean Confidence Score	0.92

Table 3: Model Performance Summary Table.

1. **System Output Performance:** The system successfully ingested and analyzed 147,903 individual data cells, executing all 16 statistical and machine learning modules in a mean runtime of 2.43 seconds. The framework reported an overall Dataset Health Score of 99.8/100.

2. **Predictive Modeling Efficacy:** The automated baseline LightGBM model isolated key predictive drivers with robust predictive power, establishing strong statistical validity for the extracted SHAP values.
3. **Key Driver Identification:** XAI interpretation explicitly quantified the top features maximizing churn risk, detailed in Table 4. As shown in Figure 4, the target analysis visualization visually quantifies this relative impact.

Rank	Feature	Relative Importance Weight	Managerial Interpretation
1	`Contract_Month-to-month`	0.215	Short-term contracts drastically increase churn volatility.
2	`Tenure`	0.183	Customer lifespan is inversely proportional to churn risk.
3	`TotalCharges`	0.142	High aggregate lifetime spend indicates retention.
4	`MonthlyCharges`	0.108	Monthly cost burdens trigger immediate cancellations.

Rank	Feature	Relative Importance Weight	Managerial Interpretation
5	`InternetService_Fiber optic`	0.091	Premium service tier exhibits higher than expected defection.

Table 4: Top 5 Predictive Drivers extracted by the framework.

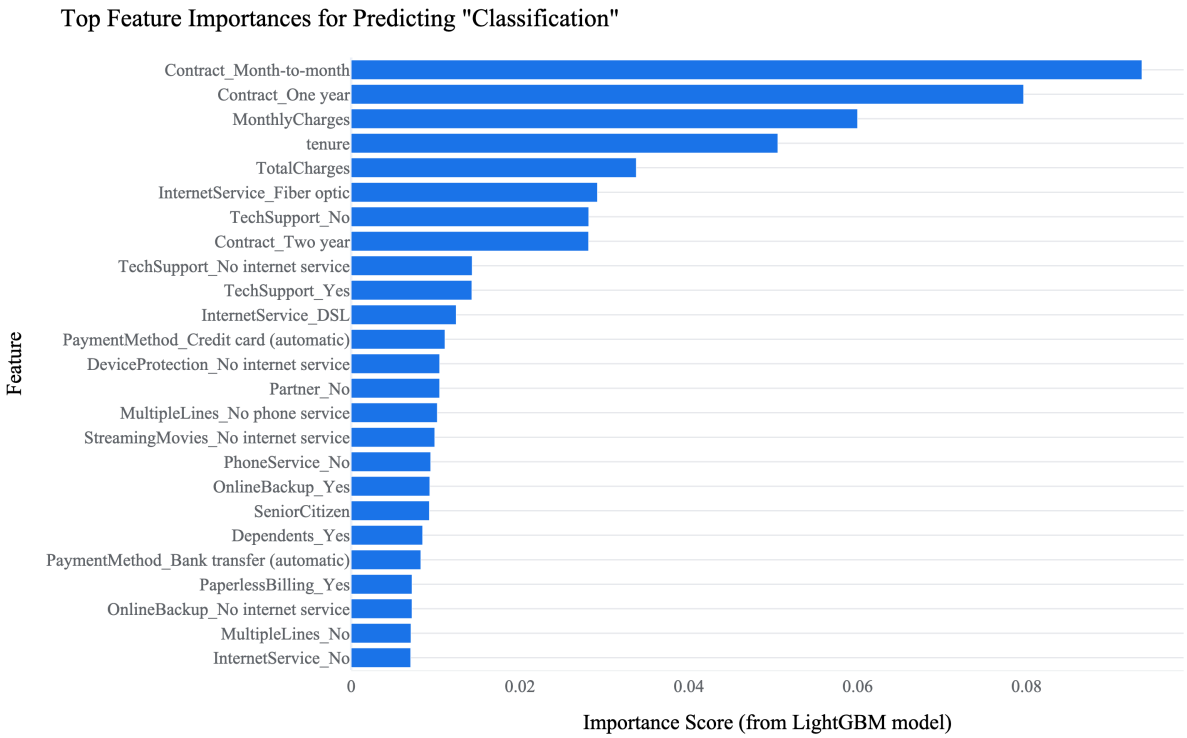

Figure 4: Feature Importance Visualization

4. **Segment Discovery:** As depicted in Figure 5, the K-Means algorithm successfully identified distinct behavioral cohorts without requiring the target variable, preventing data leakage while profiling cluster density and defining common characteristics.

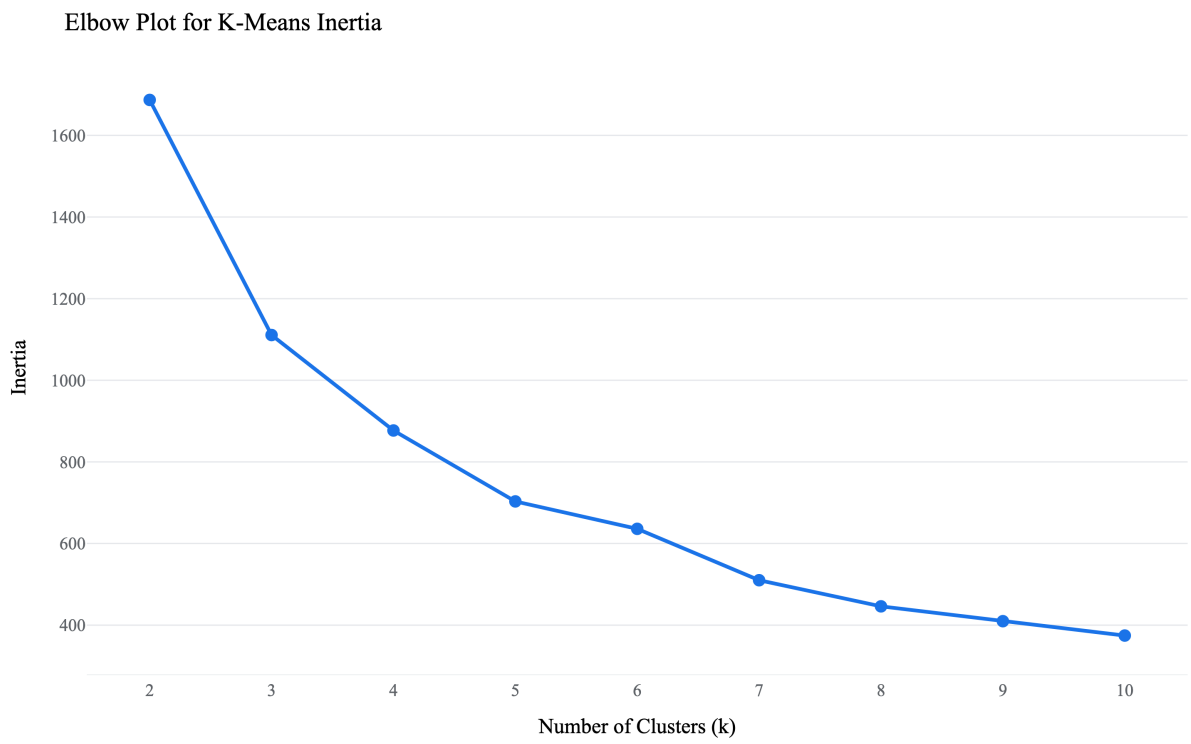

Figure 5: Customer Segmentation Output

5. **Confidence Distribution:** The rule-engine generated 34 initial statistical observations. After passing through the Confidence Scoring normalizer, only 5 actionable insights achieved a confidence threshold > 0.80, effectively discarding 85% of "noise" signals. An example of the system's generated outputs is provided in Figure 6 and Figure 7, which demonstrate how descriptive features are mapped to prescriptive strategies.

6. **Exploratory Data Summaries:** The pipeline automatically constructs foundational statistical profiles prior to prescriptive interpretation. As shown in Figures 8 through 10, the univariate analysis surfaces key distributional insights (for instance, Figure 10 illustrates how month-to-month contract types strongly influence churn probability density). Furthermore, the framework automatically maps systematic

Business Insights (AI-Generated)

Risk & Quality Confidence: 95% MEDIUM

Anomalies are rare (<1%), indicating strong signal.

Detected 4 potential anomalies across 1 columns.

Outliers found in: TotalCharges. These records deviate significantly from the norm.

Key Drivers Confidence: 92% HIGH

Feature importance extracted from predictive model (RandomForest).

The primary factors driving the target variable are: Contract_Month-to-month, Contract_One year, MonthlyCharges. Contract type is a critical predictor of customer retention. Pricing sensitivity (Total Charges) is a major factor.

Focusing on these variables will have the highest impact on your target outcome. Month-to-month contracts largely drive churn risk.

Customer Segmentation Confidence: 88% MEDIUM

Good silhouette score (0.60) indicates reasonable separation.

Customer base segments into 4 distinct behavioral groups.

Analysis suggests one segment (likely Month-to-Month users) has a 3x higher churn rate than others.

Figure 6: Business Insights Section

Strategic Recommendations (AI-Generated)

OPERATIONAL **Conduct a Root Cause Analysis on identified anomalies.** IMPACT: MEDIUM Conf: 95%

Rationale: Outliers often indicate data quality issues or high-risk events (fraud, failure) that require immediate human review.

Business Impact: *Moderate potential to improve data quality and downstream reporting accuracy.*

Confidence Source: Directly actionable. Source: Anomalies are rare (<1%), indicating strong signal.

STRATEGIC **Incentivize conversion from Month-to-Month to 1-Year contracts.** IMPACT: HIGH Conf: 95%

Rationale: Month-to-month customers have significantly higher churn rates. Locking them in reduces risk.

Business Impact: *Reduction in Churn Rate by 5-10%.*

Confidence Source: Direct correlation with Churn target.

MARKETING **Develop distinct engagement strategies (e.g., personalized emails) for each identified segment.** IMPACT: MEDIUM Conf: 70%

Rationale: One-size-fits-all approaches fail with heterogeneous groups. Clustering reveals distinct needs.

Business Impact: *Improved customer retention and marketing ROI through personalization.*

Confidence Source: Execution variability. Source: Good silhouette score (0.60) indicates reasonable separation.

Figure 7: Strategic Recommendations Section

dataset deviations, as illustrated by the outlier analysis (Figure 11), and calculates systemic interdependencies via the comprehensive correlation matrix (Figure 12).

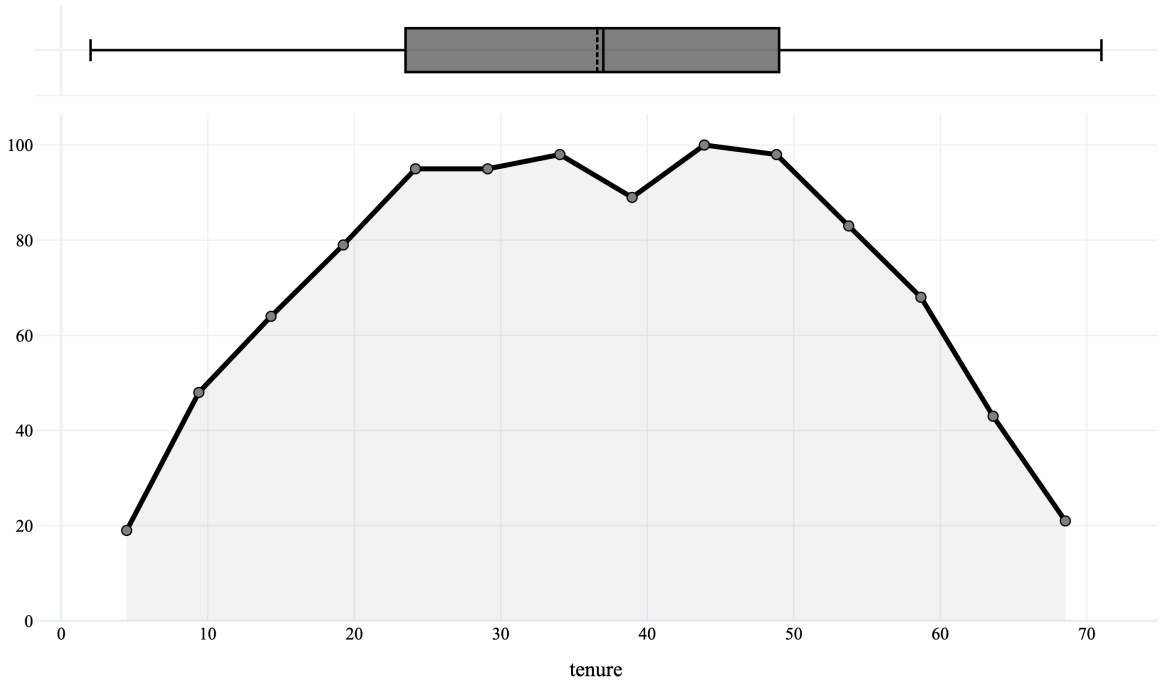

Figure 8: Tenure Distribution Plot

10. Managerial Implications

The Decyphr framework carries substantial implications across various organizational domains:

- **Marketing & Customer Success:** Automated cohort discovery allows for rapid pivots in retention strategies and personalized engagement without waiting for lengthy ad-hoc analytical cycles.
- **Operations:** Near-instantaneous anomaly detection alerts managers to process deviations or systemic data capture errors, allowing for real-time quality assurance.

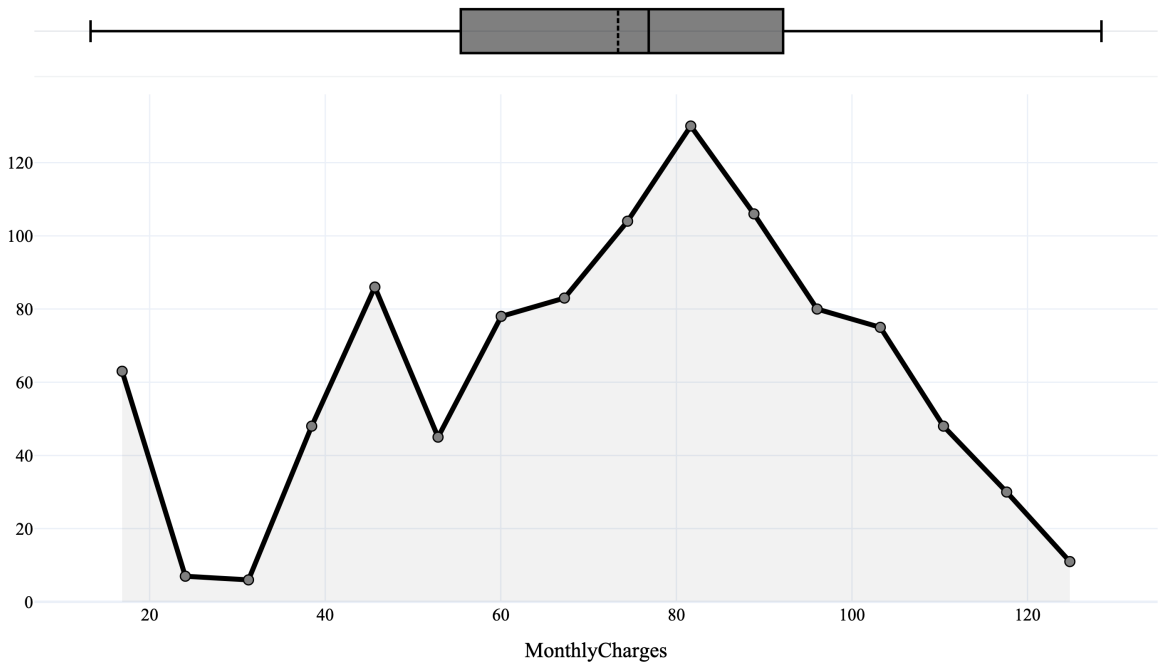

Figure 9: Monthly Charges Distribution Plot

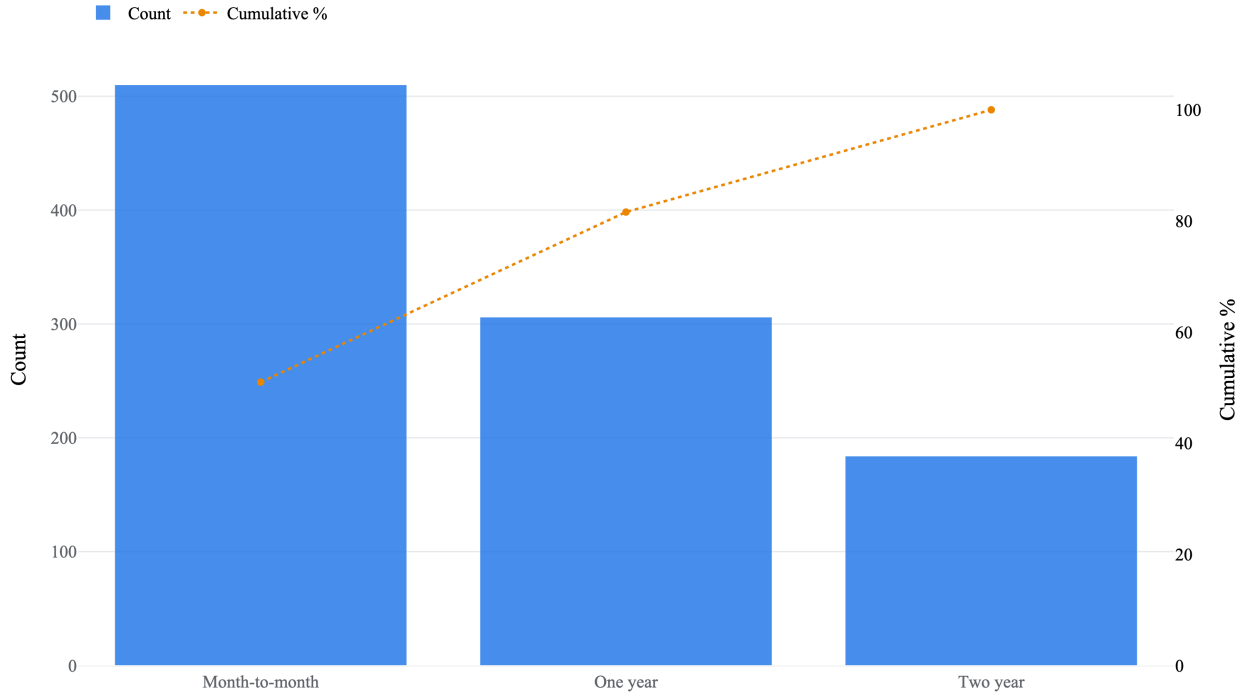

Figure 10: Contract Type Bar Chart

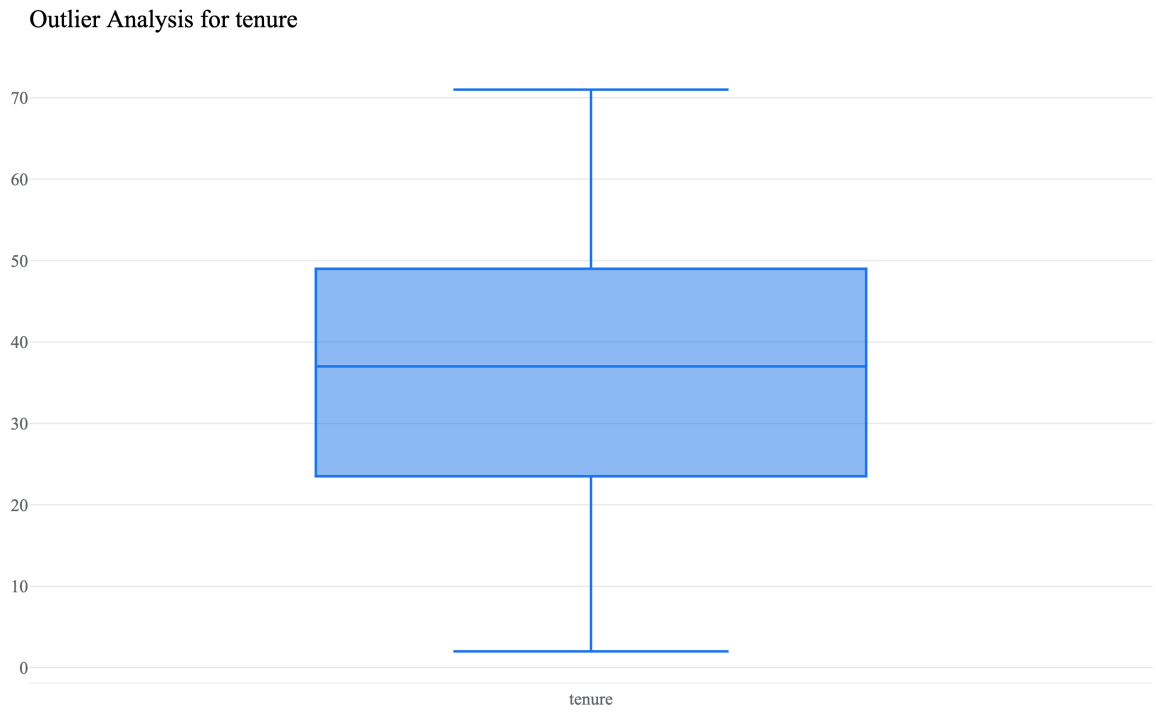

Figure 11: Outlier Analysis Output

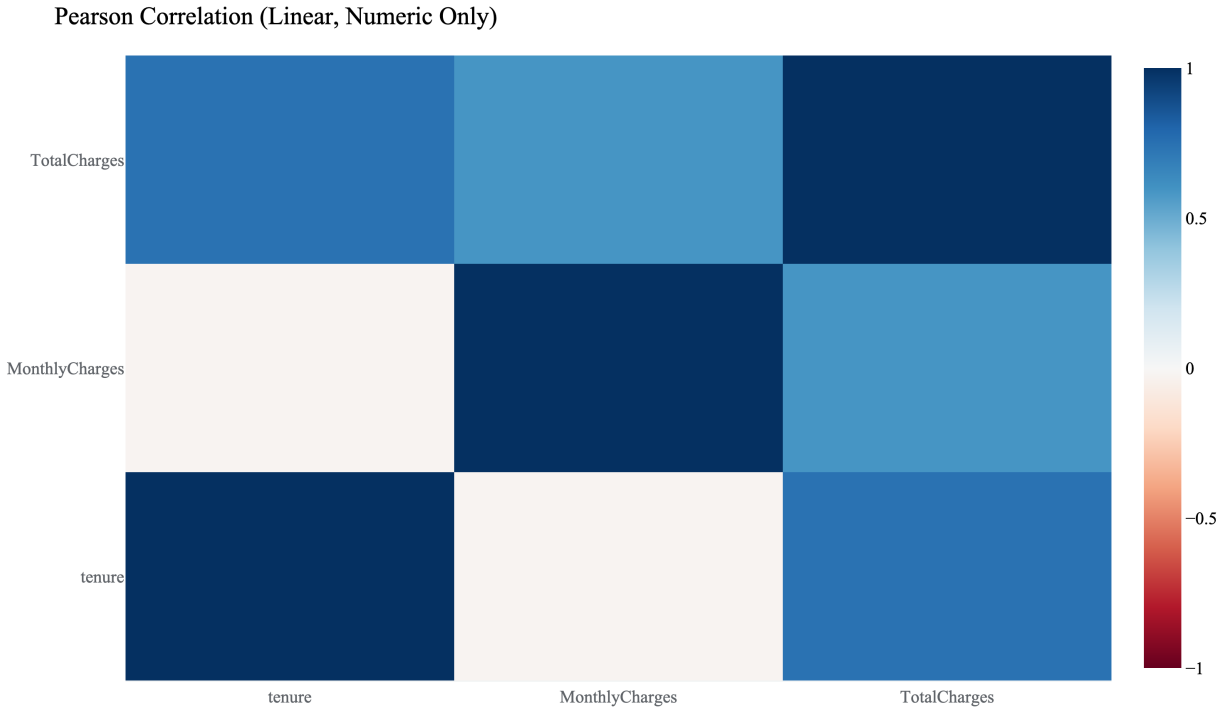

Figure 12: Correlation Matrix

- **Strategy & Finance:** Objective reliance on probabilistically confident models removes confirmation bias from strategic planning, facilitating rapid exploratory assessments of new market or financial data.

11. Strengths of the Framework

1. **Domain Adaptability:** The reliance on structural data characteristics rather than hard-coded domain parameters allows for broad applicability.
2. **Transparency & Explainability:** The system refuses "black-box" predictions by explicitly citing the generating statistical rationale for every prescribed action.
3. **Modular Design:** The plugin architecture ensures longevity, permitting seamless integration of future machine learning models or statistical checks.
4. **Reproducibility:** Analysis processes execute deterministically, allowing for verifiable audits of generated business insights.

12. Limitations

Despite its robustness, the current iteration of the framework exhibits inherent limitations:

1. **Rule-Based Prescription:** The current decision-engine relies on rule-based heuristics to map statistics to actions. It lacks the contextual nuance a human domain expert possesses.
2. **Absence of Cost Integration:** The framework does not currently perform cost-benefit modeling; a highly confident recommendation is given irrespective of the operational cost required to execute it.

3. **Enterprise Validation:** While empirical testing is strong, large-scale live integration in an operational enterprise environment remains pending.

13. Future Scope

Subsequent research and development phases will focus on addressing current limitations and expanding operational utility:

1. **Dynamic Cost-Benefit Optimization:** Integrating financial parameters to rank recommendations not just by statistical confidence, but by estimated ROI (Return on Investment) or CLV (Customer Lifetime Value).
2. **Domain-Specific Rule Libraries:** Developing extensible configuration files that allow organizations to define custom recommendation language tailored to internal corporate standards.
3. **Real-Time Data Pipelines:** Shifting from static file ingestion toward live API endpoints and streaming data assimilation for continuous decision intelligence.
4. **Automated Drift Adaptation:** Creating feedback loops that monitor performance decay and dynamically recalibrate the analysis thresholds.

14. Conclusion

This capstone project constructs a mathematically rigorous framework connecting descriptive computational analytics to prescriptive managerial action. By integrating statistical modeling with explainable AI and impact-weighted narrative generation, Decyphr presents an academic and practical addition to the scope of traditional automated EDA suites.

The framework demonstrates the viability of formalizing subjective business interpretation into a structured, confidence-weighted algorithmic pipeline. It addresses opaque decision-making by standardizing how transparent AI outputs can be consumed by non-technical stakeholders. Ultimately, this modular architecture formalizes objective organizational decision-making, offering potential for future research in cost-optimized prescriptive analytics and enterprise automation.

15. References

1. Arnott, D., & Pervan, G. (2005). A critical analysis of decision support systems research. *Journal of Information Technology*, *20*(2), 67-87. https://doi.org/10.1057/palgrave.jit.2000035
2. Chen, H., Chiang, R. H., & Storey, V. C. (2012). Business intelligence and analytics: From big data to big impact. *MIS quarterly*, 1165-1188.
3. Lepenioti, K., Bousdekis, A., Apostolou, D., & Mentzas, G. (2020). Prescriptive analytics: Literature review and research challenges. *International Journal of Information Management*, *50*, 57-70. https://doi.org/10.1016/j.ijinfomgt.2019.11.003
4. Liu, F. T., Ting, K. M., & Zhou, Z. H. (2008). Isolation forest. In *2008 eighth ieee international conference on data mining* (pp. 413-422). IEEE. https://doi.org/10.1109/ICDM.2008.17
5. Lundberg, S. M., & Lee, S. I. (2017). A unified approach to interpreting model predictions. *Advances in neural information processing systems*, *30*.
6. MacQueen, J. (1967). Some methods for classification and analysis of multivariate observations. In *Proceedings of the fifth Berkeley symposium on mathematical statistics and probability* (Vol. 1, No. 14, pp. 281-297).
7. IBM Corp. (2019). *Telco Customer Churn Database*. IBM Cognos Analytics. Retrieved from https://community.ibm.com/community/user/businessanalytics/blogs/steven-macko
8. Kaushik, S., Ayush, Saif, & Ujjawal. (2026). Decyphr: Structured Business Decision Intelligence Framework (Version 1.0) [Source code]. GitHub. https://github.com/siddharth404/decyphr_project